

The Acceleration Gap: Why Professional Services Must Move Now

Johan Strömquist | February 2026

In late December 2025, OpenAI co-founder Andrej Karpathy posted something that caught the attention of the technology world. “I’ve never felt this much behind as a programmer,” he wrote. “The profession is being dramatically refactored. I have a sense I could be 10x more powerful if I just properly string together what has become available over the last year, and a failure to claim the boost feels decidedly like a skill issue.”

This is not a junior employee venting about the pace of change. This is one of the most technically capable people on earth – someone who helped build the systems driving this shift – saying that even he feels behind relative to what is now possible. If that is the experience at the absolute frontier of AI capability, what does it mean for the rest of us?

The answer, I believe, is that we are watching a gap open in real time. A gap between organisations and individuals who are adapting to the new reality of AI-augmented work and those who are not. And unlike earlier phases of AI adoption, this gap is no longer growing at a steady, manageable pace. It is compounding.

The data on divergence

The evidence for this divergence is no longer anecdotal. It is showing up in large-scale surveys, economic research, and corporate earnings reports.

PwC’s 29th Annual Global CEO Survey, released at Davos in January 2026 and covering 4,454 CEOs across 95 countries, paints a stark picture. Only 12% of CEOs report that AI has delivered both cost savings and revenue gains. 56% report no significant financial benefit. 22% say AI has actually increased their costs.¹

The headline, however, is not that AI doesn't work. It's the divergence. The 12% achieving dual benefits are 2.6 times more likely to have embedded AI deeply across their products and services. Organisations with strong AI foundations – responsible AI frameworks, enterprise-wide integration – are three times more likely to report meaningful financial returns. The companies that are getting it right are pulling away from those that aren't, and the distance between them is growing.¹

BCG's September 2025 study "From Potential to Profit" quantifies this even more sharply. Only 5% of firms globally are what BCG calls "future-built" – generating substantial value from AI. Another 35% are scaling and beginning to see returns. The remaining 60% are, in BCG's assessment, "reaping hardly any material value" despite substantial investment. The future-built 5% achieve five times the revenue increases and three times the cost reductions of the lagging majority. They allocate 15% of their AI budgets to agents, compared with near-zero among the bottom 60%.²

The employee picture compounds the problem. Section's AI Proficiency Report, based on 5,000 knowledge workers across the US, UK, and Canada, found that less than 3% of the workforce are AI practitioners or experts who integrate AI meaningfully into their workflows. 97% are using AI poorly or not at all. Even employees who had completed their company's AI training scored only 40 out of 100 on proficiency – still classified as "experimenters" rather than practitioners.³

Meanwhile, there is a perception gap between leadership and the people doing the work. Workday's "Elevating Human Potential" study of 2,500 workers across 22 countries found that 62% of leaders report their people, processes, and technology do not work effectively together on AI.⁴ McKinsey's "Superagency" research found that over 40% of executives say AI saves them more than eight hours per week, while 40% of non-management workers report zero time savings.⁵ Leadership and the front line are living in different realities when it comes to AI.

The acceleration gap

New York Times technology columnist Kevin Roose captured the experiential dimension of this in a widely discussed post in January 2026: "I follow AI adoption pretty closely, and I have never seen such a yawning inside/outside gap. People in San Francisco are putting multi-agent Claude Swarms in charge of their lives, consulting chatbots before every decision... People elsewhere are still trying to get approval to use Copilot in Teams."⁶

Roose's observation resonated because it articulated something many people were sensing but hadn't named. The gap between AI early adopters and everyone else is no longer a gentle slope – it has become a cliff.

For much of 2023 and 2024, the difference between early adopters and the mainstream was a matter of degree. Some people used AI more than others, but the basic understanding of what AI could do – answer questions, draft text, summarise documents – was broadly shared. The gap existed, but it grew at a roughly consistent rate.

Something changed in late 2025 and early 2026. The combination of new model capabilities, autonomous agents that can execute multi-step work without supervision, and tools that put those capabilities in the hands of non-developers created what feels like a phase shift. The frontier of AI usage accelerated sharply. And because AI capabilities beget more advanced use cases, which beget more advantage, which widens the gap further, the divergence now compounds.

Linear growth in an exponential environment is a compounding disadvantage. That is the acceleration gap.

The IMF has been tracking the labour market implications of this shift at a macro level. In January 2024, IMF Managing Director Kristalina Georgieva warned that AI could affect 60% of jobs in advanced economies and 40% globally, with roughly half of exposed roles potentially benefiting from AI integration and the other half facing “lower labour demand, leading to lower wages, reduced hiring, and in the most extreme cases, job disappearance.”⁷ At Davos in January 2026, she described AI’s impact on labour markets as “a tsunami.”⁸

Research published in January 2026 by the National Bureau of Economic Research adds important granularity. Manning and Aguirre’s study of 356 occupations covering 96% of American employment introduces the concept of “adaptive capacity” – measuring not just which jobs are exposed to AI, but which workers have the financial resources, skill transferability, geographic mobility, and demographic profile to weather displacement. Their finding: 6.1 million US workers face both high AI exposure and low adaptive capacity. 86% of these most vulnerable workers are women, concentrated in administrative and clerical roles.⁹ The Brookings Institution’s accompanying analysis emphasises that these workers are “concentrated enough geographically to reach efficiently” for policy intervention – meaning the disruption is not abstract. It has a postcode.¹⁰

What actually changed

To understand why the gap is widening now, you need to understand what changed in the underlying technology – and it is a more fundamental shift than many observers appreciate.

Most people’s mental model of AI is still the chatbot: you ask a question, you get an answer. You paste in some text, the AI edits it. This is the Copilot/ChatGPT paradigm, and it is roughly where the majority of enterprises sit today.

What emerged in late 2025 and has accelerated through early 2026 is something qualitatively different: AI as an autonomous agent. Rather than answering a question, an agent receives an objective and executes it end-to-end. It reads files, writes code, searches the web, creates documents, tests its own work, and iterates – all without real-time human oversight. Multiple agents can work in parallel. They can run while you sleep.

This is not an incremental improvement on the chatbot. It is a category change. As one observer put it, thinking of AI as a chatbot is like thinking of electricity as “the thing that powers a lightbulb.”

The implications for knowledge work are profound. Nathan Lambert, an AI researcher and founder of Interconnects, captured the personal experience of this shift in his essay "Get Good at Agents": "Since trying Claude Code with Opus 4.5, my work life has shifted closer to trying to adapt to a new way of working with agents. This new style of work feels like a larger shift than the era of learning to work with chat-based AI assistants." He describes the feeling of moving "from using the power tool to pointing the army" and argues that "pointing the agents more effectively is far more useful than me spending a few more hours grinding on a problem."¹¹

The shift Lambert describes – from executor to director, from wielding the tool to orchestrating the workforce – is precisely the shift that most organisations have not yet made. And it has specific implications for how we think about skills, operating models, and commercial structures.

The skills for the new era

If execution is becoming cheap and abundant, what remains scarce?

The answer, I think, falls into two categories. The first is what we might call **agent management** – the practical skills of working with AI agents effectively. Systems design thinking: the ability to architect coherent wholes rather than implement individual components. Ambitious task scoping: knowing how to give agents meaningful end-to-end work rather than small cleanup tasks. Asynchronous work management: orchestrating multiple agents working in the background without real-time monitoring. And output validation at scale: knowing when AI output is correct without having to review every line.

The second category is what we might call **enterprise operation** – the strategic skills of knowing what to build and why. Domain expertise: understanding the specific constraints, regulations, stakeholder dynamics, and institutional knowledge of a particular field. Problem recognition: the ability to reinterpret business challenges as problems that AI-augmented approaches can solve. And process redesign: the willingness to rethink entire workflows from scratch rather than layering AI on top of existing processes.

Domain expertise deserves particular emphasis. In a world where execution is cheap, the knowledge of what to execute and whether the result is correct becomes dramatically more valuable. The senior professional who understands a regulatory landscape, a therapeutic area, a compliance regime, or a client's unstated constraints has skills that AI cannot replicate. But those skills are only fully leveraged when they are deployed as direction-setting and validation – not consumed by committee cycles and document production.

There is a tension here that organisations must confront honestly. As Lambert and others have noted, if AI handles execution, how do junior staff learn the domain expertise that currently comes from years of doing the work? If we compress the execution layer, where does the next generation of experts come from? This is not a reason to slow down. But it is a problem that must be designed for explicitly, not left to sort itself out.

The professional services reckoning

For professional services firms – consulting, advisory, legal, regulatory – these dynamics converge into a direct challenge to the business model.

The traditional model is built on a simple equation: value equals time multiplied by expertise. Clients pay for hours of senior people thinking about their problems. More hours generally signals more rigour, more thoroughness, more quality. The entire commercial structure – pricing, staffing, utilisation targets, partner economics – rests on this assumption.

AI breaks the assumption. When a task that took a team of analysts two weeks can be accomplished by an agent in two hours, the hours-based pricing model does not merely come under pressure. It becomes indefensible. Not in theory, not eventually, but now.

The data confirms that this reckoning is already underway. According to Consultancy.uk, 73% of consulting clients now prefer pricing models tied to measurable business outcomes rather than time spent.¹² McKinsey's internal generative AI tools are saving consultants approximately 30% of their time on projects.¹² ConsultingQuest reports that projects that once took ten weeks can now be executed in six, representing a 30-40% reduction in cost base, and that AI automates roughly 60% of analytical work in consulting engagements.¹³

ConsultingQuest also identifies what they call the "AI Dividend" problem: delivery costs have fallen dramatically within firms, but client fees have remained static. The efficiency gains are being "invisibly converted into profit" rather than passed to clients or reinvested in capability.¹³ This opacity is unsustainable. As clients become more aware of what AI can do – and competitors demonstrate it – the gap between production cost and billing will be exposed.

The largest firms are already restructuring in response. Deloitte announced in January 2026 that it is scrapping traditional job titles as AI reshapes its delivery model.¹⁴ BCG has grown from zero AI revenue to \$2.7 billion annually, with AI consulting now representing 20% of revenue.¹⁵ Accenture reported \$4.1 billion in generative AI revenue for fiscal year 2025 and a 27% increase in project efficiency after deploying AI internally.¹⁶ McKinsey expects 40% of its business to be AI-related in the near future.¹⁵

Perhaps most telling is what is happening to the traditional consulting pyramid. McKinsey, BCG, and Bain have frozen starting salaries for the third consecutive year.¹⁷ PwC projects that tax and assurance associate hiring will fall by 32% between 2025 and 2028, with audit hiring declining by 39%.¹⁸ One emerging model compresses junior layers by 50-70%, with AI operators replacing much of the entry-level workforce.¹⁹ HBR notes that AI is "dismantling the traditional hiring model" of professional services firms.²⁰

And the disruption is not only coming from within the industry. OpenAI has launched a consulting division with a reported \$10 million minimum engagement, embedding "Forward-Deployed

Engineers" directly with enterprise clients – cutting out traditional consulting intermediaries entirely.²¹

FirmWise estimates that at least 20% of all professional services firms will have to fundamentally change their revenue model in the next five years to accommodate AI.²²

The asymmetry of inaction

For any professional services leader reading this, the temptation is to wait. To let the technology mature, the hype subside, the dust settle. To treat AI as a 2027 strategy problem rather than a 2026 operational one.

I think this is precisely wrong, and I think the reason it is wrong is that the costs of action and inaction are deeply asymmetric.

If you invest in AI transformation and AI turns out to be overhyped, you have spent time and money upskilling your people. They are more capable. You have modernised some processes. The cost is recoverable.

If you do not invest and AI turns out to be as transformative as the evidence suggests, you are structurally uncompetitive. Your pricing model is indefensible against AI-augmented competitors. Your best people leave for firms that give them better tools. You lose mandates you used to win. The cost is potentially existential.

The acceleration gap is real. It is supported by data from PwC, BCG, McKinsey, the IMF, the NBER, and the lived experience of thousands of professionals at the frontier. The gap compounds. And the window for moving from the wrong side to the right side narrows with every quarter that passes.

The good news is that the skills that matter most in the AI era – domain expertise, strategic judgment, the ability to recognise and frame the right problems – are precisely the skills that experienced professionals already have. The shift is not about replacing those skills. It is about redeploying them: away from the production of deliverables and toward the direction-setting and validation that only human expertise can provide.

The question is not whether this shift is coming. It is whether we choose to lead it or be led by it.

References

1. PwC, "29th Annual Global CEO Survey: Leading through uncertainty in the age of AI," January 19, 2026. 4,454 CEOs across 95 countries, surveyed September-November 2025. Full report (PDF) | Press release ↵
2. BCG, "From Potential to Profit: Closing the AI Impact Gap," September 2025. Report | Press release ↵
3. Section, "The AI Proficiency Report," 2025. Survey of 5,000 knowledge workers from 1,000+ person companies in the US, UK, and Canada. Report ↵
4. Workday, "Elevating Human Potential: The AI Skills Revolution," 2025. Global study of 2,500 full-time workers across 22 countries, fielded by Hanover Research, November-December 2024. Report ↵
5. McKinsey, "Superagency in the workplace: Empowering people to unlock AI's full potential at work," 2025. Report ↵
6. Kevin Roose, X (formerly Twitter) post, January 25, 2026. Discussed in Gizmodo and Scott Loftess, "The Reality Gap" ↵
7. IMF Staff Discussion Note SDN/2024/001, "Gen-AI: Artificial Intelligence and the Future of Work," January 14, 2024. Authors: Cazzaniga, Jaumotte, Li, Melina, Panton, Pizzinelli, Rockall, Tavares. Publication | Georgieva blog post ↵
8. Georgieva remarks at Davos, January 2026. Reported in Fortune. Further remarks at World Government Summit, Dubai, February 2026: IMF ↵
9. Manning, S.J. and Aguirre, T., "How Adaptable Are American Workers to AI-Induced Job Displacement?" NBER Working Paper No. 34705, January 2026. NBER ↵
10. Brookings Institution, "Measuring US workers' capacity to adapt to AI-driven job displacement," January 21, 2026. Authors: Manning, Aguirre, Muro, Methkupally. Analysis ↵
11. Nathan Lambert, "Get Good at Agents," Interconnects AI (Substack), January 2026. Essay ↵
12. Consultancy.uk, "Is AI about to kill the billable hour consulting model?" 2025. Article ↵
13. ConsultingQuest, "AI Impact on Consulting Economics and Value Sharing," 2025. Analysis ↵
14. Fortune, "Deloitte scraps traditional job titles as AI reshapes Big 4 consulting firms," January 22, 2026. Article ↵
15. Medium / Takafumi Endo, "How AI is Redefining Strategy Consulting: McKinsey, BCG, and Bain," 2025. Article. See also BrainForge, "How Big Consulting Firms Profit from AI Consulting." Article ↵

16. Accenture Newsroom, "OpenAI and Accenture accelerate enterprise reinvention with advanced AI," December 2025. Press release ↵
17. Irish Times, "Top consultancies freeze starting salaries as AI threatens pyramid model," December 1, 2025. Article ↵
18. Management Consulted, "PwC entry-level recruiting to slow," 2025. Article ↵
19. FourWeekMBA, "AI: The Pyramid Compression in the Consulting Business Model." Analysis ↵
20. Harvard Business Review, "How AI Is Upending How Consulting Firms Hire Talent," October 2025. Article. See also Fortune, "Harvard Business Review Survey: Only 6% of companies trust AI agents," December 9, 2025. Article ↵
21. The Information, "OpenAI Takes a Page from Palantir, Doubles Down on Consulting Services." Article. See also The Finance Story, "OpenAI Enters AI Consulting, Challenging Big 4." Article ↵
22. FirmWise, "AI in Professional Services: 2025 State of the Industry Report." Report ↵